Siamese Network:

A Siamese Network is a neural network architecture specifically designed for similarity learning, where the goal is to determine whether two inputs are similar or dissimilar. It consists of two identical subnetworks that process two inputs independently and produce feature embeddings, which are then compared using a distance metric (e.g., Euclidean distance) to quantify similarity. Siamese Networks are widely used for image similarity search because they are explicitly trained to learn a similarity function rather than classifying inputs. Unlike traditional methods like SIFT, which relies on handcrafted feature detection and matching, Siamese Networks learn features directly from the data, making them more adaptable to specific tasks. Compared to CNN-based methods like ResNet, which focus on extracting hierarchical features for classification tasks, Siamese Networks are trained to optimize for pairwise similarity, making them more suitable for one-shot or few-shot learning tasks where labeled data is scarce.

The advantages of Siamese Networks include their ability to generalize well to unseen data with minimal training samples (e.g., in one-shot learning), their flexibility in learning task-specific similarity metrics, and their adaptability to complex, high-level features that are difficult for handcrafted methods like SIFT to capture. Additionally, Siamese Networks are robust to variations in scale, rotation, and other transformations due to the learned embeddings. However, there are a few limitations. Siamese Networks often require large amounts of carefully constructed training pairs to learn effective embeddings, and their performance heavily depends on the quality and diversity of the training data. They are computationally expensive compared to SIFT and may require significant resources for training. Furthermore, if not trained properly, the network may fail to generalize, leading to suboptimal performance. Despite these challenges, Siamese Networks are highly effective for applications requiring image similarity, face recognition, and visual search, especially when task-specific adaptations are needed.

Implementation of Siamese Network:

1. `load_data()`

- **Data Normalization**: The images are divided by `255.0` to scale pixel intensities from their original range (0–255) to a range of `0–1`. This improves model convergence during training.
- **Adding a Channel Dimension**: The line `[..., np.newaxis]` reshapes the dataset from `(N, H, W)` to `(N, H, W, 1)` to include a channel dimension, as convolutional layers expect input with this format. Here, `N` is the number of samples, `H` is the height, and `W` is the width.

This preprocessing ensures the dataset is ready for input into a convolutional neural network.

2. `create_triplets(x, y, num_triplets=10000)`

This function generates triplets for training (anchor, positive, and negative images). It works as follows:

- `num_triplets=10000`: Specifies the number of triplets to generate.
- `digit_indices = [np.where(y == i)[0] for i in range(num_classes)]`: Groups the dataset indices by class (`num_classes=10` for Fashion MNIST).
- **Anchor and Positive Selection**: An anchor and a positive image are randomly sampled from the same class.
- **Negative Selection**: A negative image is sampled from a different class. The function uses a loop to ensure the negative label differs from the anchor label. The generated triplets are used to train the Siamese network, ensuring it learns to differentiate between similar and dissimilar images.

3. `build_siamese_network(input_shape)`

This function creates the base feature extraction network for the Siamese architecture:

- `input_shape`: The shape of the input images (e.g., `(28, 28, 1)` for Fashion MNIST).
- **Convolutional Layers**:
  - `Conv2D(64, (3, 3), activation='relu')`: Applies 64 filters of size $3x3$ to extract features. The ReLU activation introduces non-linearity.
  - `Conv2D(128, (3, 3), activation='relu')`: Uses 128 filters for deeper feature extraction.
- **Pooling Layers**:
  - `MaxPooling2D()`: Reduces the spatial dimensions of the feature maps, retaining the most important information while down-sampling.
- **Dense Layers**:
  - `Dense(128, activation='relu')` and `Dense(64, activation='relu')`: Fully connected layers compress the feature maps into dense embeddings.
- **Dropout (`Dropout(0.2)`)**: Randomly drops 20% of the neurons during training to prevent overfitting.
- **BatchNormalization**: Normalizes the output of the dense layers to stabilize and speed up training.

This feature extractor outputs a `64-dimensional` embedding for each input image.

4. `triplet_loss(y_true, y_pred, margin=0.2)`

This function calculates the **triplet loss**, ensuring that embeddings for similar images (anchor and positive) are closer than embeddings for dissimilar images (anchor and negative) by at least a margin.

- **Inputs**:
  - `y_true`: True labels (not used in the loss calculation).
  - `y_pred`: A tensor containing the embeddings for anchor, positive, and negative images.
- **Embedding Distances**:
  - `positive_distance`: Squared Euclidean distance between anchor and positive embeddings.
  - `negative_distance`: Squared Euclidean distance between anchor and negative embeddings.
  - These are computed as `tf.reduce_sum(tf.square(anchor - positive), axis=-1)`.
- **`margin=0.2`**: Specifies the minimum distance margin required between positive and negative embeddings.
- **Output**:
  - The loss is calculated as `tf.reduce_mean(tf.maximum(positive_distance - negative_distance + margin, 0))`. It ensures the loss is 0 when the margin is satisfied.

5. `call(self, inputs)`

This function is part of a custom layer (`TripletStack`) and is used to stack the anchor, positive, and negative embeddings together into a single tensor. The stacked embeddings are then used to calculate the triplet loss during training.

- **`TripletStack`**: This is a subclass of TensorFlow's `Layer` class.
- **`inputs`**: A tuple containing the embeddings for the anchor, positive, and negative images.
- **`tf.stack([...], axis=1)`**: Combines the embeddings into a single tensor by stacking them along the first axis. This results in a tensor of shape `[batch_size, 3, embedding_size]`, where:
  - `3` corresponds to the anchor, positive, and negative embeddings.

6. `create_triplet_model(input_shape)`

This function creates the full triplet model by combining the base Siamese network with three inputs:

- **Inputs**:
  - `input_anchor`, `input_positive`, and `input_negative` represent the anchor, positive, and negative images.
- **Base Model**:
  - The `base_model` (created using `build_siamese_network`) is used to extract embeddings for all three inputs.
- **Stacking Embeddings**:
  - The embeddings are stacked using the `TripletStack` layer, and the model outputs these embeddings for triplet loss calculation.

7. `train_model(model, triplets_train, triplets_test, epochs=15, batch_size=64)`

This function trains the triplet model:

- **epochs=15**: Specifies the number of passes over the entire dataset during training.
- **batch_size=64**: Determines how many triplets are processed at once. Larger batch sizes can improve training stability but require more memory.
- **Training Loop**:
  - The model is trained on `triplets_train` and validated on `triplets_test` using the triplet loss.

8. `visualize_embeddings(base_model, x, y)`

This function visualizes the embeddings generated by the base model:

- **Embedding Extraction**: The embeddings are computed using `base_model.predict(x)`.
- **t-SNE (`TSNE(n_components=2)`)**:
  - Reduces the dimensionality of embeddings to 2D for visualization.
- **Plotting**:
  - Each embedding is plotted using a scatter plot, with colors representing different classes.

9. `calculate_distance(embedding_a, embedding_b)`

- **Implementation**: `np.linalg.norm(embedding_a - embedding_b, axis=-1)` calculates the distance between embeddings

10. `generate_top_matches(base_model, x_test, y_test, query_index, top_k=5)`

This function retrieves the top-k most similar images to a query:

- `query_index`: The index of the query image.
- `top_k=5`: The number of similar images to return.
- **Steps**:
  - Compute embeddings for all test images using the `base_model`.
  - Calculate distances between the query embedding and all other embeddings.
  - Sort the distances and retrieve the indices of the `top_k` smallest distances.

11. `evaluate_similarity(base_model, x_test, y_test, top_k=5, num_queries=100)`

This function evaluates the similarity performance of the model:

- `top_k=5`: Evaluates how often the correct class is among the top-5 most similar embeddings.
- `num_queries=100`: Limits the evaluation to 100 random queries for efficiency.

12. `plot_similar_images(query_image, database_images, query_embedding, database_embeddings, top_k=5, cmap='viridis')`

This function visualizes the query image and the top-k most similar images:

- `top_k=5`: Displays the top-5 most similar images.
- `cmap='viridis'`: Uses the viridis colormap for visualization.
- **Steps**:
  - Compute distances between the query embedding and the database embeddings.
  - Sort distances and retrieve indices for the top-k similar images.
  - Plot the query image alongside the top-k similar images.

Summary of the Siamese Network Results

1. t-SNE Visualization of Embeddings

- **Embeddings Distribution**:
  - The t-SNE plot shows a 2D representation of the high-dimensional embeddings generated by the Siamese model.
  - Each cluster corresponds to a specific class (as denoted by different colors in the legend).
- **Observation**:
  - The embeddings are well-separated, meaning the model has successfully learned to map similar images closer together and dissimilar images further apart in the embedding space.
  - This clear separation indicates strong feature extraction capabilities by the Siamese network.

2. Performance Metrics

- **Precision, Recall, and Accuracy**:
  - Precision: **0.8180**
  - Recall: **0.0206**
  - Accuracy: **0.9300**
- **Observation**:
  - The high **accuracy** of 93% shows that the model performs well in correctly classifying or matching images.

3. Top-5 Similar Images Retrieval

- **Query and Retrieved Images**:
  - The query image (on the left) represents pants.
  - The top-5 similar images retrieved from the database also represent pants, showcasing the network's ability to identify visual similarity.
- **Observation**:
  - The retrieved images are visually consistent with the query, confirming the model's effectiveness in similarity-based retrieval.

Overall Summary

1. **Training**:
   - The network trained well, with consistent reduction in loss values over epochs.
   - No signs of overfitting, as indicated by the small gap between training and validation loss.
2. **Embeddings**:

- ○ The t-SNE visualization demonstrates the network's ability to cluster similar images effectively in the embedding space.
3. **Performance**:
   - ○ High precision and accuracy indicate the network's reliability in image similarity search tasks.
4. **Retrieval**:
   - ○ The network successfully retrieves semantically similar images from the database.